

Pre-Process

```
In [16]: import pandas as pd
import numpy as np
df = pd.read_csv('../data/compas-scores-two-years.csv')
df.columns
```

```
Out[16]: Index(['id', 'name', 'first', 'last', 'compas_screening_date', 'sex', 'dob',
'age', 'age_cat', 'race', 'juv_fel_count', 'decile_score',
'juv_misd_count', 'juv_other_count', 'priors_count',
'days_b_screening_arrest', 'c_jail_in', 'c_jail_out', 'c_case_number',
'c_offense_date', 'c_arrest_date', 'c_days_from_compas',
'c_charge_degree', 'c_charge_desc', 'is_recid', 'r_case_number',
'r_charge_degree', 'r_days_from_arrest', 'r_offense_date',
'r_charge_desc', 'r_jail_in', 'r_jail_out', 'violent_recid',
'is_violent_recid', 'vr_case_number', 'vr_charge_degree',
'vr_offense_date', 'vr_charge_desc', 'type_of_assessment',
'decile_score.1', 'score_text', 'screening_date',
'v_type_of_assessment', 'v_decile_score', 'v_score_text',
'v_screening_date', 'in_custody', 'out_custody', 'priors_count.1',
'start', 'end', 'event', 'two_year_recid'],
dtype='object')
```

```
In [17]: compas_df = df
races = ['African-American', 'Caucasian']

compas_df=compas_df[compas_df.race.isin(races)]

compas_df.head()

#use only specific columns for future prediction
compas_df_filtered = compas_df.loc[:, ['race']].join(compas_df.select_dtypes(include=['int64']))
#drop id and duplicate columns
compas_df_filtered = compas_df_filtered.drop(['id', 'decile_score.1', 'priors_count.1'], axis=1)
compas_df_filtered.head()
```

Out[17]:

	race	age	juv_fel_count	decile_score	juv_misd_count	juv_other_count	priors_count	is_recid	is_violent_recid	v_decile_score	start	end
1	African-American	34	0	3	0	0	0	1	1	1	9	0
2	African-American	24	0	4	0	1	4	1	0	3	0	0
3	African-American	23	0	8	1	0	1	0	0	6	0	1
6	Caucasian	41	0	6	0	0	14	1	0	2	5	0
8	Caucasian	39	0	1	0	0	0	0	0	1	2	0

In [18]: `compas_df_filtered["race"].replace(['African-American', 'Caucasian'],[0, 1], inplace=True)`

In [19]:

```

y = compas_df_filtered["two_year_recid"]
X = compas_df_filtered.drop(["two_year_recid"], axis=1)
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test=train_test_split(X,y,test_size=0.2, random_state=50, stratify=y)
print(X_train.shape, y_train.shape, X_test.shape, y_test.shape)
X_train, X_valid, y_train, y_valid=train_test_split(X_train,y_train,test_size=0.2, random_state=67, stratify=y_train)
print(X_train.shape, y_train.shape, X_valid.shape, y_valid.shape)

(4920, 13) (4920,) (1230, 13) (1230,)
(3936, 13) (3936,) (984, 13) (984,)

```

Baseline

In [20]:

```

from sklearn.linear_model import LogisticRegression
clf = LogisticRegression().fit(X_train, y_train)
clf.score(X_test,y_test)

```

```
C:\Users\zhang\Anaconda3\lib\site-packages\sklearn\linear_model\_logistic.py:818: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

```
Increase the number of iterations (max_iter) or scale the data as shown in:  
https://scikit-learn.org/stable/modules/preprocessing.html  
Please also refer to the documentation for alternative solver options:  
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression  
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG,  
0.9829268292682927
```

Out[20]:

```
In [21]: index_b= (X_test["race"]==0)  
X_test_b=X_test[index_b]  
y_test_b=y_test[index_b]  
X_test_w=X_test[-index_b]  
y_test_w=y_test[-index_b]
```

```
In [22]: index_b= (X_train["race"]==0)  
X_train_b=X_train[index_b]  
y_train_b=y_train[index_b]  
X_train_w=X_train[-index_b]  
y_train_w=y_train[-index_b]
```

```
In [24]: #calibration  
print(clf.score(X_test_b,y_test_b))  
print(clf.score(X_test_w,y_test_w))  
print("calibration score: ",abs(clf.score(X_test_b,y_test_b)-clf.score(X_test_w,y_test_w)))  
  
0.9807692307692307  
0.9860557768924303  
calibration score: 0.005286546123199565
```

A5

```
In [25]: import torch as t  
import torch.nn as nn
```

Process the data for the model

```
In [26]: data_filtered_afram = compas_df_filtered[compas_df_filtered['race']==0]
data_filtered_cau = compas_df_filtered[compas_df_filtered['race']==1]

#X_afram = data_filtered_afram[["race", "age", "priors_count", "juv_count", 'score_text', "sex", "c_charge_degree", "decile_sc
#y_afram = data_filtered_afram[['two_year_recid']]

X_afram = data_filtered_afram.drop(["two_year_recid"], axis=1)
y_afram = data_filtered_afram[["two_year_recid"]]

X_cau = data_filtered_cau.drop(["two_year_recid"], axis=1)
y_cau = data_filtered_cau[["two_year_recid"]]

#X_cau = data_filtered_cau[["race", "age", "priors_count", "juv_count", 'score_text', "sex", "c_charge_degree", "decile_score"
#y_cau = data_filtered_cau[['two_year_recid']]

In [27]: X_train_afram, X_test_afram, y_train_afram, y_test_afram = train_test_split(X_afram, y_afram, test_size=0.2)
X_train_cau, X_test_cau, y_train_cau, y_test_cau = train_test_split(X_cau, y_cau, test_size=0.2)
```

Use tensor flow the modify the dataset

```
In [28]: def tensorX(X):
        return tf.tensor(np.array(X)).to(tf.float32)

        def tensorY(X,Y):
            return tf.from_numpy(np.array(Y).astype('float32')).reshape(X.shape[0], 1)

In [29]: X_train_afram = tensorX(X_train_afram)
X_test_afram = tensorX(X_test_afram)
X_train_cau = tensorX(X_train_cau)
X_test_cau = tensorX(X_test_cau)

y_train_afram = tensorY(X_train_afram,y_train_afram)
y_test_afram = tensorY(X_test_afram,y_test_afram)
y_train_cau = tensorY(X_train_cau,y_train_cau)
y_test_cau = tensorY(X_test_cau,y_test_cau)
```

Algorithm: Prejudice Remover Regularizer Loss Function

```
In [30]: #update LogisticRegression function
class LogisticRegression(nn.Module):
```

```

def __init__(self,data):
    super(LogisticRegression, self).__init__()
    self.w = nn.Linear(data.shape[1], out_features=1, bias=True)
    self.sigmod = nn.Sigmoid()
def forward(self, x):
    w = self.w(x)
    output = self.sigmod(w)
    return output

```

```

In [31]: #Define the Prejudice Remover Regularizer Loss Function
class PRLoss():
    def __init__(self, eta=1.0):
        super(PRLoss, self).__init__()
        self.eta = eta

    def PI(self,output_afram,output_cau):
        N_afram = t.tensor(output_afram.shape[0])
        N_cau = t.tensor(output_cau.shape[0])
        # calculate P[y/s]
        P_ys = t.stack((t.sum(output_afram),t.sum(output_cau)),axis=0) / t.stack((N_afram,N_cau),axis=0)
        # calculate P[y]
        P_y = t.sum(t.cat((output_afram,output_cau),0)) / (X_train_afram.shape[0]+X_train_cau.shape[0])
        # calculate PI
        PI_s1y1 = output_afram * (t.log(P_ys[1]) - t.log(P_y))
        PI_s1y0 = (1- output_afram) * (t.log(1-P_ys[1]) - t.log(1-P_y))
        PI_s0y1 = output_cau * (t.log(P_ys[0]) - t.log(P_y))
        PI_s0y0 = (1- output_cau) * (t.log(1-P_ys[0]) - t.log(1-P_y))
        PI = t.sum(PI_s1y1) + t.sum(PI_s1y0) + t.sum(PI_s0y1) + t.sum(PI_s0y0)
        PI = self.eta * PI
        return PI

```

Algorithm:Logistical regression with Prejudice Remover Regularizer

```

In [32]: #This function has a hyperparameter eta which is the size of the regulating term in the loss function.
#The fit method will give us the accuracy and calibration of the fitted model
class LogisticRegressionWithPRR():

    def __init__(self, eta=0.0, epochs=100, lr = 0.01):
        super(LogisticRegressionWithPRR, self).__init__()
        self.eta = eta
        self.epochs = epochs
        self.lr = lr

```

```

def fit(self,X_train_afram,y_train_afram,X_train_cau,y_train_cau,X_test_afram, y_test_afram, X_test_cau, y_test_cau
    #LogisticRegression model
    model_afram = LogisticRegression(X_train_afram)
    model_cau = LogisticRegression(X_train_cau)

    criterion = nn.BCELoss()
    PI = PRLoss(eta=self.eta)
    epochs = self.epochs

    #L2 regularization (non-zero weight_decay)
    optimizer = t.optim.Adam(list(model_afram.parameters())+ list(model_cau.parameters()), self.lr, weight_decay=1e

    for epoch in range(epochs):
        #train
        model_afram.train()
        model_cau.train()

        #zero out the gradients
        optimizer.zero_grad()

        #compute loss
        output_afram = model_afram(X_train_afram)
        output_cau = model_cau(X_train_cau)
        log_loss = criterion(output_afram, y_train_afram)+ criterion(output_cau, y_train_cau)
        PI_loss = PI.PI(output_afram,output_cau)
        loss = PI_loss +log_loss

        loss.backward()
        optimizer.step()
    #eval
    model_afram.eval()
    model_cau.eval()

    #calculate accuracy
    #Accuracy is the average of correctly predicted labels of two groups
    #Calibration is the difference of accuracy bewtween the two groups
    y_pred_afram = (model_afram(X_test_afram) >= 0.5)
    y_pred_cau = (model_cau(X_test_cau) >= 0.5)

    #sum of correct prediction/total num
    accuracy_afram = t.sum(y_pred_afram == y_test_afram) / y_test_afram.shape[0]
    accuracy_cau = t.sum(y_pred_cau == y_test_cau) / y_test_cau.shape[0]

    accuracy = (accuracy_afram + accuracy_cau) / 2
    calibration = t.abs(accuracy_afram - accuracy_cau)

```

```
return accuracy.item(), calibration.item(), accuracy_afram.item() , accuracy_cau.item()
```

```
In [41]: PR = LogisticRegressionWithPRR(eta = 860, epochs = 100, lr = 1e-04)
PR.fit(X_train_afram,y_train_afram,X_train_cau,y_train_cau, X_test_afram, y_test_afram, X_test_cau, y_test_cau)
```

```
Out[41]: (0.43259480595588684,
0.15643197298049927,
0.5108107924461365,
0.3543788194656372)
```

Tuning Hyperparameter eta

```
In [42]: #Tuning the hyperparameter to find the best model
eta=np.linspace(0, 100, num=1000)
hist_acc=np.zeros(1000)
hist_cal=np.zeros(1000)
hist_afram = np.zeros(1000)
hist_cau = np.zeros(1000)
for i in range(1000):
    PR = LogisticRegressionWithPRR(eta = eta[i], epochs = 100, lr = 1e-04)
    results = PR.fit(X_train_afram,y_train_afram,X_train_cau,y_train_cau, X_test_afram, y_test_afram, X_test_cau, y_test_cau)
    hist_acc[i]=results[0]
    hist_cal[i]=results[1]
    hist_afram[i]=results[2]
    hist_cau[i]=results[3]
```

```
In [43]: clf_accuracy = np.max(hist_acc)
print("Accuracy:")
np.max(hist_acc)
```

```
Accuracy:
Out[43]: 0.7574337720870972
```

```
In [44]: A5_cal_score = np.min(hist_cal)
print("Calibration Score:")
hist_cal[np.argmax(hist_acc)] #calibration decreases
```

```
Calibration Score:
Out[44]: 0.0851324200630188
```

```
In [45]: #clf_accuracy = hist_acc[np.argmin(hist_cal)]  
print("Accuracy:")  
hist_acc[np.argmin(hist_cal)]
```

Accuracy:
Out[45]: 0.6058223247528076

```
In [46]: #A5_cal_score = np.min(hist_cal)  
print("Calibration Score:")  
np.min(hist_cal) #minimum calibration
```

Calibration Score:
Out[46]: 0.0018687844276428223

```
In [47]: #np.sort(hist_cal)[0:5]  
hist_cal_new = hist_cal  
np.sort(hist_cal_new)[0:5]
```

Out[47]: array([0.00186878, 0.00220457, 0.00247699, 0.00727418, 0.00727418])

```
In [55]: np.where(hist_cal <= 0.0025)
```

Out[55]: (array([158, 745, 897], dtype=int64),)

```
In [56]: hist_acc[158]
```

Out[56]: 0.6203831434249878

```
In [57]: A5_acc = hist_acc[158]  
A5_calibration = np.sort(hist_cal_new)[2]
```

A6

```
In [96]: import numpy as np  
import pandas as pd  
from sklearn.model_selection import train_test_split  
from sklearn.linear_model import LogisticRegression  
from sklearn.metrics import confusion_matrix  
import matplotlib.pyplot as plt
```


Data Loading and Preprocessing

```
In [97]: compas_df= pd.read_csv('../data/compas-scores-two-years.csv')

#Use the data of African American and Caucasian

races = ['African-American','Caucasian']

compas_df=compas_df[compas_df.race.isin(races)]

compas_df.head()

#use only specific columns for future prediction
compas_df_filtered = compas_df.loc[:, ['race']].join(compas_df.select_dtypes(include=['int64']))
#drop id and duplicate columns
compas_df_filtered = compas_df_filtered.drop(['id','decile_score.1', 'priors_count.1'], axis=1)
#compas_df_filtered = df[["two_year_recid", "race", "age", "priors_count", "juv_count", 'score_text', "sex", "c_charge_degree"]
#compas_df_filtered.head()
```

Implemtnaion of algorithms

```
In [98]: #for the algorithms below, we apply on the whole dataset
X=compas_df_filtered
s=X.race
e=X.decile_score
y=X.two_year_recid

def PARTITION(X, e):
    unique_values = e.unique() # find unique values of e
    groups = {} # initialize empty dictionary to store resulting DataFrames

    for val in unique_values:
        group_df = X[e == val] # filter rows in X where e equals current unique value
        groups[val] = group_df # store resulting DataFrame in dictionary with key = unique value

    return groups # return dictionary of resulting DataFrames
```

```
In [99]: resulting_groups = PARTITION(X, e)
first_group_df = resulting_groups[1]
#print(first_group_df.head)

unique_values = e.unique()
print(unique_values)
```

```
[ 3  4  8  6  1 10  5  9  2  7]
```

```
In [100... def delta(X, gender):
    # calculate the baseline probabilities
    p_pos_base = X.two_year_recid.mean()

    # calculate probabilities conditioned on gender
    p_pos_gender = X[X.race == gender].two_year_recid.mean()

    # calculate the number of gender people in X
    G = len(X[X.race == gender])

    # calculate delta

    delta_val = G * abs(p_pos_gender - p_pos_base)

    return int(delta_val)
```

```
In [101... for group in resulting_groups.values():
    print(delta(group, 'Caucasian'))
```

```
11
10
3
1
5
4
3
1
1
2
```

Algorithm: Local massaging

In [102...

```
#Algorithm 1: Local massaging
resulting_groups = PARTITION(X, e)
updated_groups_1 = []
for group in resulting_groups.values():
    males = group[group.race == 'Caucasian']
    females = group[group.race == 'African-American']

    # Learn a ranker H for this group
    H = LogisticRegression(random_state=42,max_iter=500)
    H.fit(group.drop(['race','two_year_recid'], axis=1), group.two_year_recid)

    # rank and relabel males
    males_females = pd.concat([males,females])
    males_females['proba'] = H.predict_proba(males_females.drop(['two_year_recid','race'], axis=1))[:, 1]

    males_females = males_females.sort_values(by='proba')

    delta_males = int(delta(group, 'Caucasian'))

    males_to_relabel = males_females[(males_females['race'] == 'Caucasian') & (males_females['two_year_recid'] == 0) &
    group.loc[males_to_relabel.index, 'two_year_recid'] = 1 - males_to_relabel.two_year_recid

    # rank and relabel females
    delta_females = int(delta(group, 'African-American'))
    females_to_relabel = males_females[(males_females['race'] == 'African-American') & (males_females['two_year_recid']
    group.loc[females_to_relabel.index, 'two_year_recid'] = 1 - females_to_relabel.two_year_recid

    updated_groups_1.append(group)
```

```
C:\Users\zhang\Anaconda3\lib\site-packages\pandas\core\indexing.py:1773: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
self._setitem_single_column(ilocs[0], value, pi)
```

```
C:\Users\zhang\Anaconda3\lib\site-packages\pandas\core\indexing.py:1773: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
self._setitem_single_column(ilocs[0], value, pi)
```

```
C:\Users\zhang\Anaconda3\lib\site-packages\pandas\core\indexing.py:1773: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
self._setitem_single_column(ilocs[0], value, pi)
```

```
C:\Users\zhang\Anaconda3\lib\site-packages\pandas\core\indexing.py:1773: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
self._setitem_single_column(ilocs[0], value, pi)
```

```
C:\Users\zhang\Anaconda3\lib\site-packages\pandas\core\indexing.py:1773: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
self._setitem_single_column(ilocs[0], value, pi)
```

```
C:\Users\zhang\Anaconda3\lib\site-packages\pandas\core\indexing.py:1773: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
self._setitem_single_column(ilocs[0], value, pi)
```

```
C:\Users\zhang\Anaconda3\lib\site-packages\pandas\core\indexing.py:1773: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
self._setitem_single_column(ilocs[0], value, pi)
```

C:\Users\zhang\Anaconda3\lib\site-packages\pandas\core\indexing.py:1773: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
self._setitem_single_column(ilocs[0], value, pi)
```

C:\Users\zhang\Anaconda3\lib\site-packages\pandas\core\indexing.py:1773: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
self._setitem_single_column(ilocs[0], value, pi)
```

C:\Users\zhang\Anaconda3\lib\site-packages\pandas\core\indexing.py:1773: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
self._setitem_single_column(ilocs[0], value, pi)
```

In [103...

```
updated_df_1 = pd.concat(updated_groups_1)
updated_df_1 = updated_df_1.sort_index()
print(sum(updated_df_1['two_year_recid']!=compas_df_filtered['two_year_recid']))
#67 labels have been changed
```

67

Algorithm: Local preferential sampling

In [104...

```
#Algorithm 2: Local preferential sampling
resulting_groups = PARTITION(X, e)
updated_groups_2 = []
for group in resulting_groups.values():
    males = group[group.race == 'Caucasian']
    females = group[group.race == 'African-American']

# Learn a ranker H for this group
H = LogisticRegression(random_state=42,max_iter=500)
```

```

H.fit(group.drop(['race', 'two_year_recid'], axis=1), group.two_year_recid)

# rank delete, and relabel males
males_females = pd.concat([males, females])
males_females['proba'] = H.predict_proba(males_females.drop(['two_year_recid', 'race'], axis=1))[:, 1]
males_females = males_females.sort_values(by='proba')
delta_males = int(0.5*delta(group, 'Caucasian'))

males_to_delete = males_females[(males_females['race'] == 'Caucasian') & (males_females['two_year_recid'] == 1) & (
males_to_duplicate = males_females[(males_females['race'] == 'Caucasian') & (males_females['two_year_recid'] == 0)
group = group[~group.index.isin(males_to_delete.index)]
group = pd.concat([group, males_to_duplicate], ignore_index=True)

# rank delete, and relabel females
delta_females = int(0.5*delta(group, 'African-American'))
females_to_delete = males_females[(males_females['race'] == 'African-American') & (males_females['two_year_recid']
females_to_duplicate = males_females[(males_females['race'] == 'African-American') & (males_females['two_year_recid'
group = group[~group.index.isin(females_to_delete.index)]
group = pd.concat([group, females_to_duplicate], ignore_index=True)

updated_groups_2.append(group)

updated_df_2 = pd.concat(updated_groups_2)

#Now we have 3 data frames.compas_df_filtered is the original one. updated_df_1 is the one of algorithm1 updated_df_1 i

```

In [105...

```

# split the data into train, validation and test set

# Split data into training and validation/test sets
train_df, val_test_df = train_test_split(compas_df_filtered, test_size=0.2, random_state=42)

# Split validation/test set into validation and test sets
val_df, test_df = train_test_split(val_test_df, test_size=0.5, random_state=42)

# Split data into training and validation/test sets
train_df_algo1, val_test_df_algo1 = train_test_split(updated_df_1, test_size=0.2, random_state=42)

# Split validation/test set into validation and test sets
val_df_algo1, test_df_algo1 = train_test_split(val_test_df_algo1, test_size=0.5, random_state=42)

```

```
# Split data into training and validation/test sets
train_df_algo2, val_test_df_algo2 = train_test_split(updated_df_2, test_size=0.2, random_state=42)

# Split validation/test set into validation and test sets
val_df_algo2, test_df_algo2 = train_test_split(val_test_df_algo2, test_size=0.5, random_state=42)
```

In [106...

```
X_train1 = train_df.drop(['is_recid', 'race'], axis=1)
y_train1 = train_df['is_recid']

X_train2 = train_df_algo1.drop(['is_recid', 'race'], axis=1)
y_train2 = train_df_algo1['is_recid']

X_train3 = train_df_algo2.drop(['is_recid', 'race', 'proba'], axis=1)
y_train3 = train_df_algo2['is_recid']

# Fit logistic regression model for X_train1
logreg1 = LogisticRegression()
logreg1.fit(X_train1, y_train1)

# Fit logistic regression model for X_train2
logreg2 = LogisticRegression()
logreg2.fit(X_train2, y_train2)

# Fit logistic regression model for X_train3
logreg3 = LogisticRegression()
logreg3.fit(X_train3, y_train3)
```

```
C:\Users\zhang\Anaconda3\lib\site-packages\sklearn\linear_model\_logistic.py:818: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG,

```
C:\Users\zhang\Anaconda3\lib\site-packages\sklearn\linear_model\_logistic.py:818: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG,

```
C:\Users\zhang\Anaconda3\lib\site-packages\sklearn\linear_model\_logistic.py:818: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG,

```
LogisticRegression()
```

Out[106]:

In [107...]

```
#Get results for validation set
# For X_val1
X_val1 = val_df.drop(['is_recid', 'race'], axis=1)
y_val1 = val_df['is_recid']
y_pred_val1 = logreg1.predict(X_val1)

# For X_val2
X_val2 = val_df_algo1.drop(['is_recid', 'race'], axis=1)
y_val2 = val_df_algo1['is_recid']
y_pred_val2 = logreg2.predict(X_val2)

# For X_val3
X_val3 = val_df_algo2.drop(['is_recid', 'race', 'proba'], axis=1)
```



```
y_val3 = val_df_algo2['is_recid']
y_pred_val3 = logreg3.predict(X_val3)
```

```
In [108... #Get results for test set
# For X_test1
X_test1 = test_df.drop(['is_recid', 'race'], axis=1)
y_test1 = test_df['is_recid']
y_pred_test1 = logreg1.predict(X_test1)

# For X_test2
X_test2 = test_df_algo1.drop(['is_recid', 'race'], axis=1)
y_test2 = test_df_algo1['is_recid']
y_pred_test2 = logreg2.predict(X_test2)

# For X_test3
X_test3 = test_df_algo2.drop(['is_recid', 'race', 'proba'], axis=1)
y_test3 = test_df_algo2['is_recid']
y_pred_test3 = logreg3.predict(X_test3)
```

```
In [109... # Make a copy of the original test dataframes
test_df1 = test_df.copy()
test_df2 = test_df_algo1.copy()
test_df3 = test_df_algo2.copy()

test_df1["y_pred"] = y_pred_test1
test_df2["y_pred"] = y_pred_test2
test_df3["y_pred"] = y_pred_test3
```

```
In [110... # Define function to compute D_all and D_bad

def compute_D(test_df, y_pred):
    # Compute D_all
    D_all = test_df.loc[test_df['race'] == 'African-American', y_pred].mean() - \
            test_df.loc[test_df['race'] == 'Caucasian', y_pred].mean()

    # Compute D_bad
    decile_scores = list(range(1, 6))
    D_bad = 0
```

```

    for score in decile_scores:
        P_score = test_df.loc[(test_df['race'] == 'African-American') & (test_df['decile_score'] == score), y_pred].mean()
        P_score -= test_df.loc[(test_df['race'] == 'Caucasian') & (test_df['decile_score'] == score), y_pred].mean()
        P_y = test_df.loc[test_df['decile_score'] == score, y_pred].mean()
        D_bad += ((P_score) * (P_y))

    return D_all, D_bad

# Compute D_all and D_bad for the three logistic regression models on the three test sets
D_all_1, D_bad_1 = compute_D(test_df1, 'two_year_recid')
D_all_2, D_bad_2 = compute_D(test_df2, 'two_year_recid')
D_all_3, D_bad_3 = compute_D(test_df3, 'two_year_recid')

```

In [111...

```

# Calculate D_all and D_bad for each model on the test set
D_all_test = [D_all_1, D_all_2, D_all_3]
D_bad_test = [D_bad_1, D_bad_2, D_bad_3]

plt.plot(['Baseline', 'Local massaging', 'Local preferential sampling'], D_all_test, label='D_all')
plt.plot(['Baseline', 'Local massaging', 'Local preferential sampling'], D_bad_test, label='D_bad')


# Add axis labels and title
plt.xlabel('Model')
plt.ylabel('D')
plt.title('D_all and D_bad for 3 models on test set')

# Add Legend
plt.legend()

# Show the plot
plt.show()

print(D_all_test)
print(D_bad_test)

```



```
[0.13573754620665374, 0.09924529132195037, 0.11945715462632062]
[0.11922600741704245, 0.07281751731787592, 0.02948165449549718]
```

Comparison plots:

In [112...

```
from sklearn.metrics import accuracy_score
import matplotlib.pyplot as plt

# Define function to compute calibration
def compute_calibration(df):
    african_american = df[df['race'] == 'African-American']
    caucasian = df[df['race'] == 'Caucasian']
    african_american_accuracy = accuracy_score(african_american['two_year_recid'], african_american['y_pred'])
    caucasian_accuracy = accuracy_score(caucasian['two_year_recid'], caucasian['y_pred'])
    return abs(african_american_accuracy - caucasian_accuracy)
```

```

# Compute accuracy and calibration for test set 1
accuracy_test1 = accuracy_score(test_df1['two_year_recid'], test_df1['y_pred'])
calibration_test1 = compute_calibration(test_df1)
baseline_cal = calibration_test1 * len((e.unique()))

# Compute accuracy and calibration for test set 2
accuracy_test2 = accuracy_score(test_df2['two_year_recid'], test_df2['y_pred'])
calibration_test2 = compute_calibration(test_df2)

# Compute accuracy and calibration for test set 3
accuracy_test3 = accuracy_score(test_df3['two_year_recid'], test_df3['y_pred'])
calibration_test3 = compute_calibration(test_df3)

A5_acc = hist_acc[158]
A5_calibration = np.sort(hist_cal_new)[2]

# Plot results
accuracies = [accuracy_test1, A5_acc, accuracy_test2, accuracy_test3]
calibrations = [baseline_cal, A5_calibration, calibration_test2, calibration_test3]

models = ['Baseline', 'A5', 'A6 LM', 'A6 LPS']
accs = accuracies
cals = calibrations

fig, axs = plt.subplots(nrows=1, ncols=2, figsize=(10, 5))
axs[0].bar(models, accuracies)
axs[0].set_title('Accuracy on test sets')
axs[1].bar(models, calibrations)
axs[1].set_title('Calibration on test sets')
plt.show()

```

